

Natural Language Processing

Lecture 14:

Machine Learning: Feed-forward Neural Networks,
Autoencoders/embeddings, Dense networks


12 /11/2021

COMS W4705
Yassine Benajiba


Perceptron Expressiveness

- Simple perceptron learning algorithm, starts with an arbitrary hyperplane and adjusts it using the training data.
 - Step function is not differentiable, so no closed-form solution.
- Perceptron produces a linear separator.
 - Can only learn linearly separable patterns.
- Can represent boolean functions like **and**, **or**, **not** but not the **xor** function.


The problem with *xor*



(a) x_1 **and** x_2




(b) x_1 **or** x_2



(c) x_1 **xor** x_2



Multi-Layer Neural Networks



- Basic idea: represent any (non-linear) function as a composition of soft-threshold functions. This is a form of non-linear regression.
- Lippmann 1987: Two hidden layers suffice to represent any arbitrary region (provided enough neurons), even discontinuous functions!

Activation Functions

- One problem with perceptrons is that the **threshold function (step function)** is undifferentiable.
- It is therefore unsuitable for gradient descent.
- One alternative is the **sigmoid (logistic) function**:




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g(z) = 0 \text{ if } z \rightarrow -\infty$$


$$g(z) = 1 \text{ if } z \rightarrow \infty$$

Activation Functions

- Two other popular activation functions:




$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$



$$\text{relu}(z) = \max(z, 0)$$


Output Representation

- Many NLP Problems are multi-class classification problems.
- Each output neuron represents one class. Predict the class with the highest activation.



Softmax

- We often want the activation at the output layer to represent probabilities.
- Normalize activation of each output unit by the sum of all output activations (as in log-linear models).




$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^k \exp(z_j)}$$

The network computes a probability

$$P(c_i | \mathbf{x}; \mathbf{w})$$

Softmax

- We often want the activation at the output layer to represent probabilities.
- Normalize activation of each output unit by the sum of all output activations (as in log-linear models).



$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^k \exp(z_j)}$$

The network computes a probability

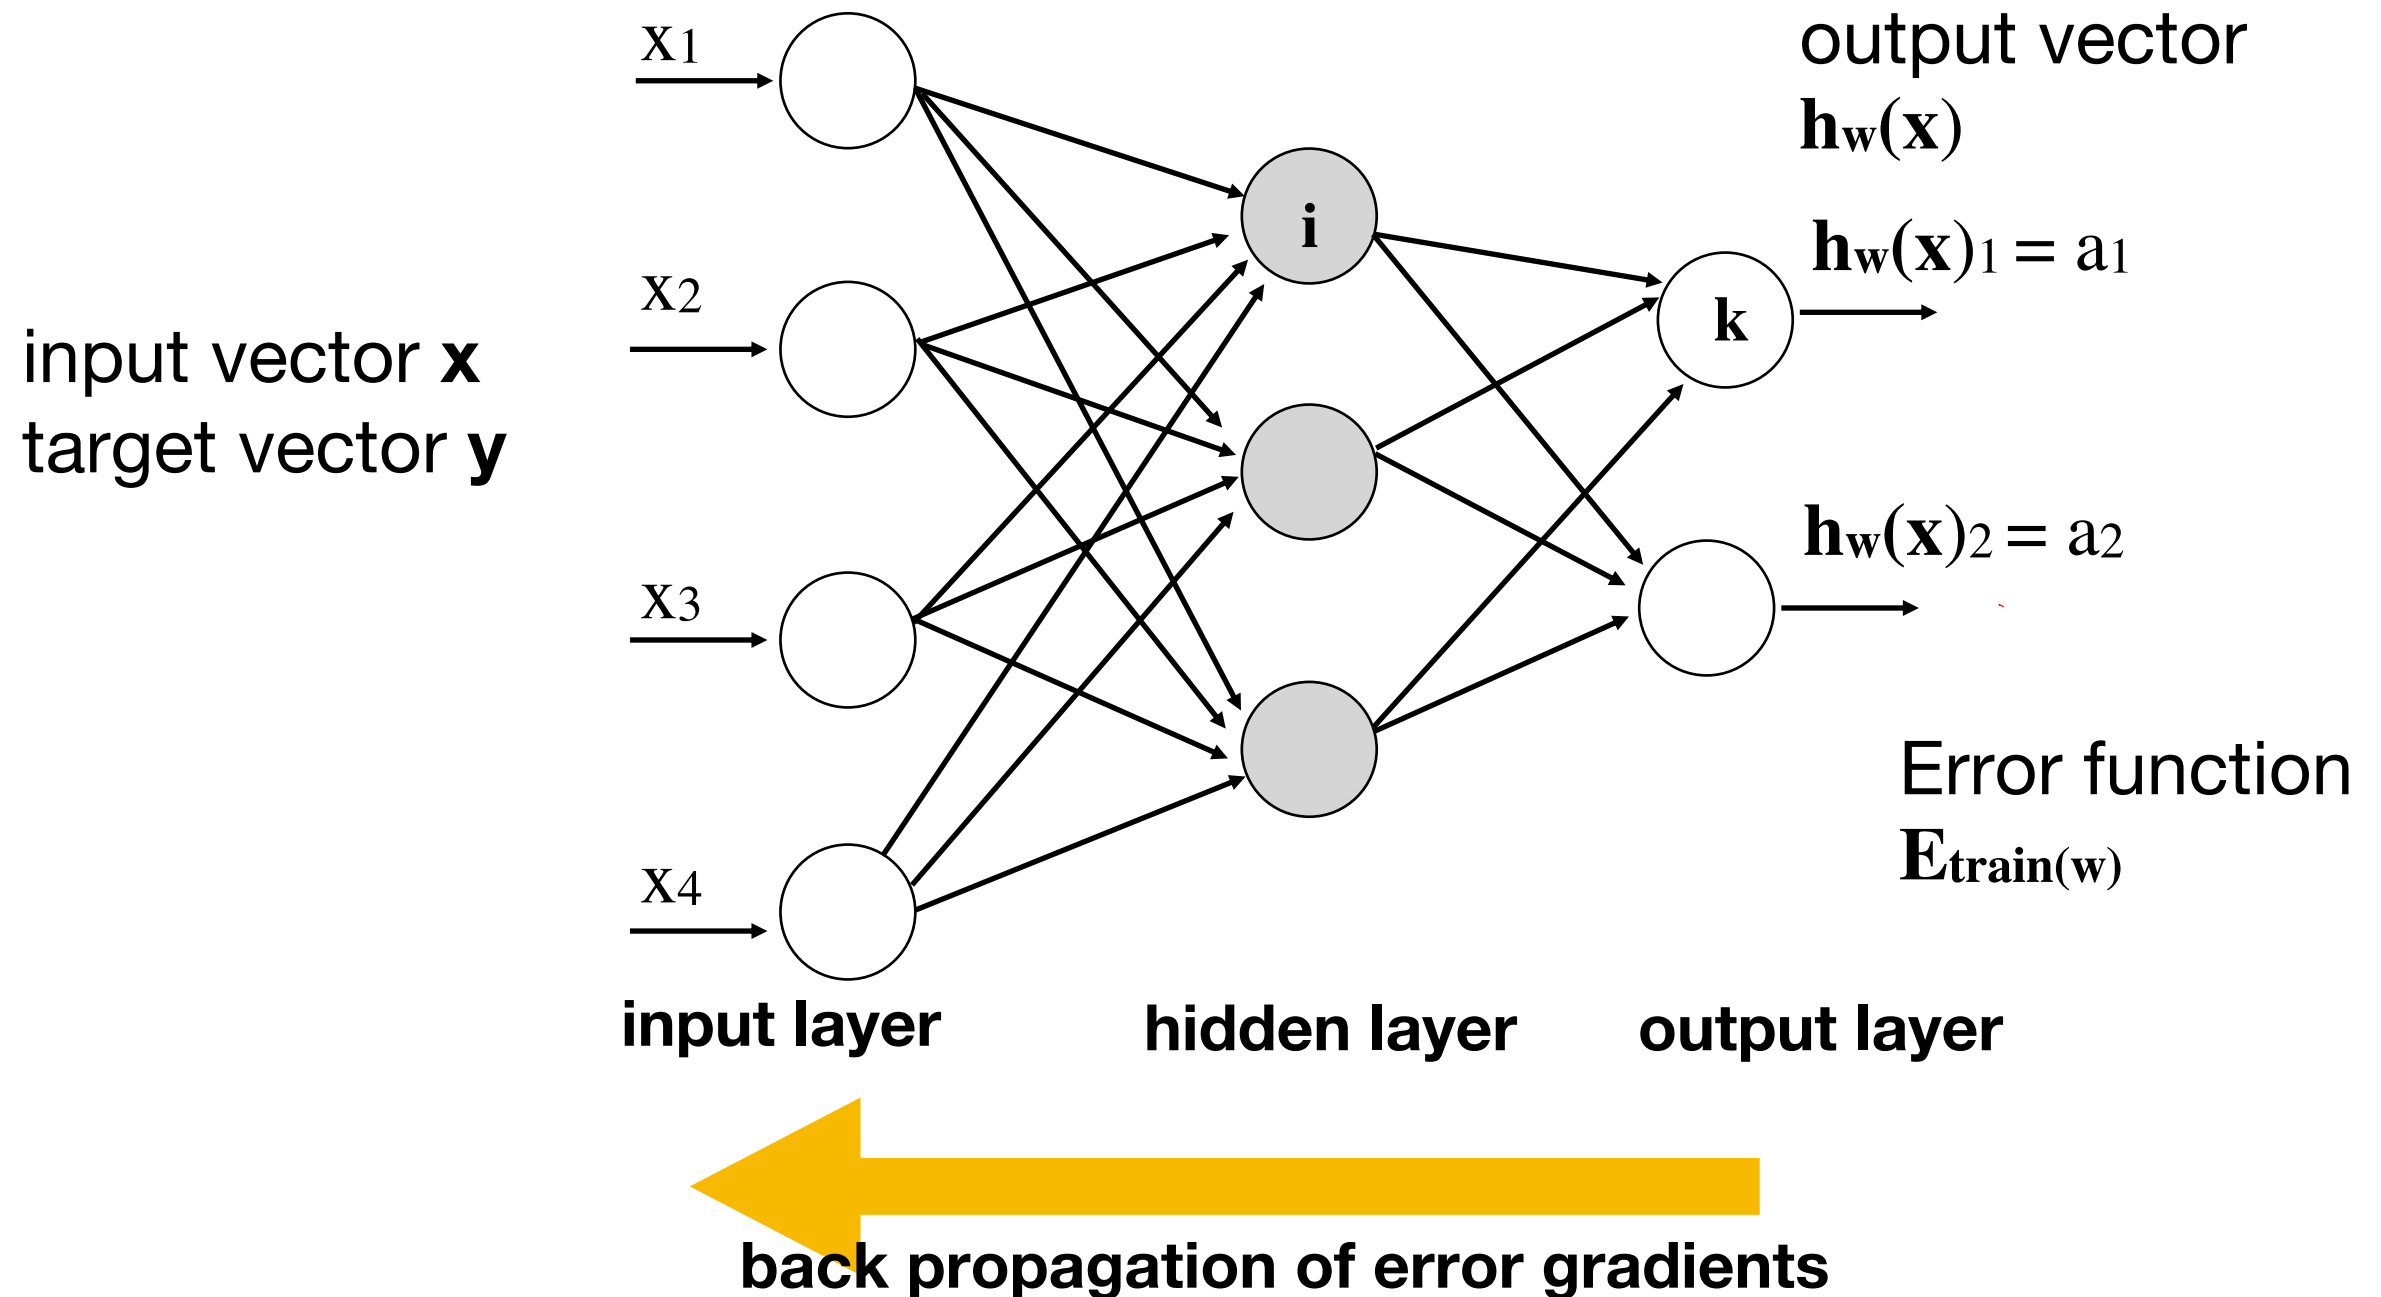
$$P(c_i | \mathbf{x}; \mathbf{w})$$

Learning in Multi-Layer Neural Networks

- Network structure is fixed, but we want to train the weights. Assume **feed-forward** neural networks: no connections that are loops.
- **Backpropagation Algorithm:**
 - Given current weights, get network output and compute loss function (assume multiple outputs / a vector of outputs).
 - Can use gradient descent to update weights and minimize loss.
 - Problem: We only know how to do this for the last layer!
 - Idea: Propagate error backwards through the network.

Backpropagation

feed-forward computation of network outputs



Negative Log-Likelihood

(also known as cross-entropy)

- Assume target output is a one-hot vector and $c(y)$ is the target class for target \mathbf{y} .
- Compute the negative log-likelihood for a single example

$$Loss(\mathbf{y}, h_{\mathbf{w}}(x)) = -\log P(c(\mathbf{y})|\mathbf{x}; \mathbf{w})$$

- Empirical error for the entire training data:

$$E_{train}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N -\log P(c(\mathbf{y}^{(i)})|\mathbf{x}^i; \mathbf{w})$$

Stochastic Gradient Descent (for a single unit)

- Goal: Learn parameters that minimize the empirical error.

Randomly initialize w

for a set number of iterations T :

shuffle training data $\mathcal{D} = (x^{(j)}, y^{(j)})|_{j=1}^n$

for $j = 1 \dots N$:

for each w_i (all weights in the network):

$$w_i \leftarrow w_i - \eta \frac{\partial}{\partial w_i} \text{Loss}(y^{(j)}, h_{\mathbf{w}}(x^{(j)}))$$

- η is the learning rate.
- It often makes sense to compute the gradient over batches of examples, instead of just one ("mini-batch").

Backpropagation

- Simplified multi-layer case (a single unit per layer):



- Stochastic Gradient Descent should perform the following update:

$$w_2 \leftarrow w_2 - \eta \frac{\partial \text{Loss}(y, f(g(x)))}{\partial w_2}$$

$$w_1 \leftarrow w_1 - \eta \frac{\partial \text{Loss}(y, f(g(x)))}{\partial w_1}$$

- Problem: How do we compute the gradient for parameters w_1 and w_2 ?

Chain Rule of Calculus

- To compute gradients for hidden units, we need to apply the chain rule of calculus:

The derivative of $f(g(x))$ is

$$\frac{df(g(x))}{dx} = \frac{df(g(x))}{dg(x)} \cdot \frac{dg(x)}{dx}$$

Backpropagation




$$\frac{\partial Loss}{\partial w_2} = \left(\frac{\partial Loss}{\partial g(f(x))} \right) \left(\frac{\partial g(f(x))}{\partial w_2} \right)$$

$$\frac{\partial Loss}{\partial w_1} = \left(\frac{\partial Loss}{\partial f(x)} \right) \left(\frac{\partial f(x)}{\partial w_1} \right)$$


$$= \left(\frac{\partial Loss}{\partial g(f(x))} \right) \left(\frac{\partial g(f(x))}{\partial f(x)} \right) \left(\frac{\partial f(x)}{\partial w_1} \right)$$

Backpropagation

forward



backward



Assume we know

$$\frac{\partial Loss}{\partial f(x)}$$

We want to compute

$$\frac{\partial Loss}{\partial x}$$

to propagate it back.


and

$$\frac{\partial Loss}{\partial w}$$


(for the weight update)

Backpropagation

forward



backward



$$\frac{\partial \text{Loss}}{\partial x} = \left(\frac{\partial \text{Loss}}{\partial f(x)} \right) \left(\frac{\partial f(x)}{\partial x} \right)$$

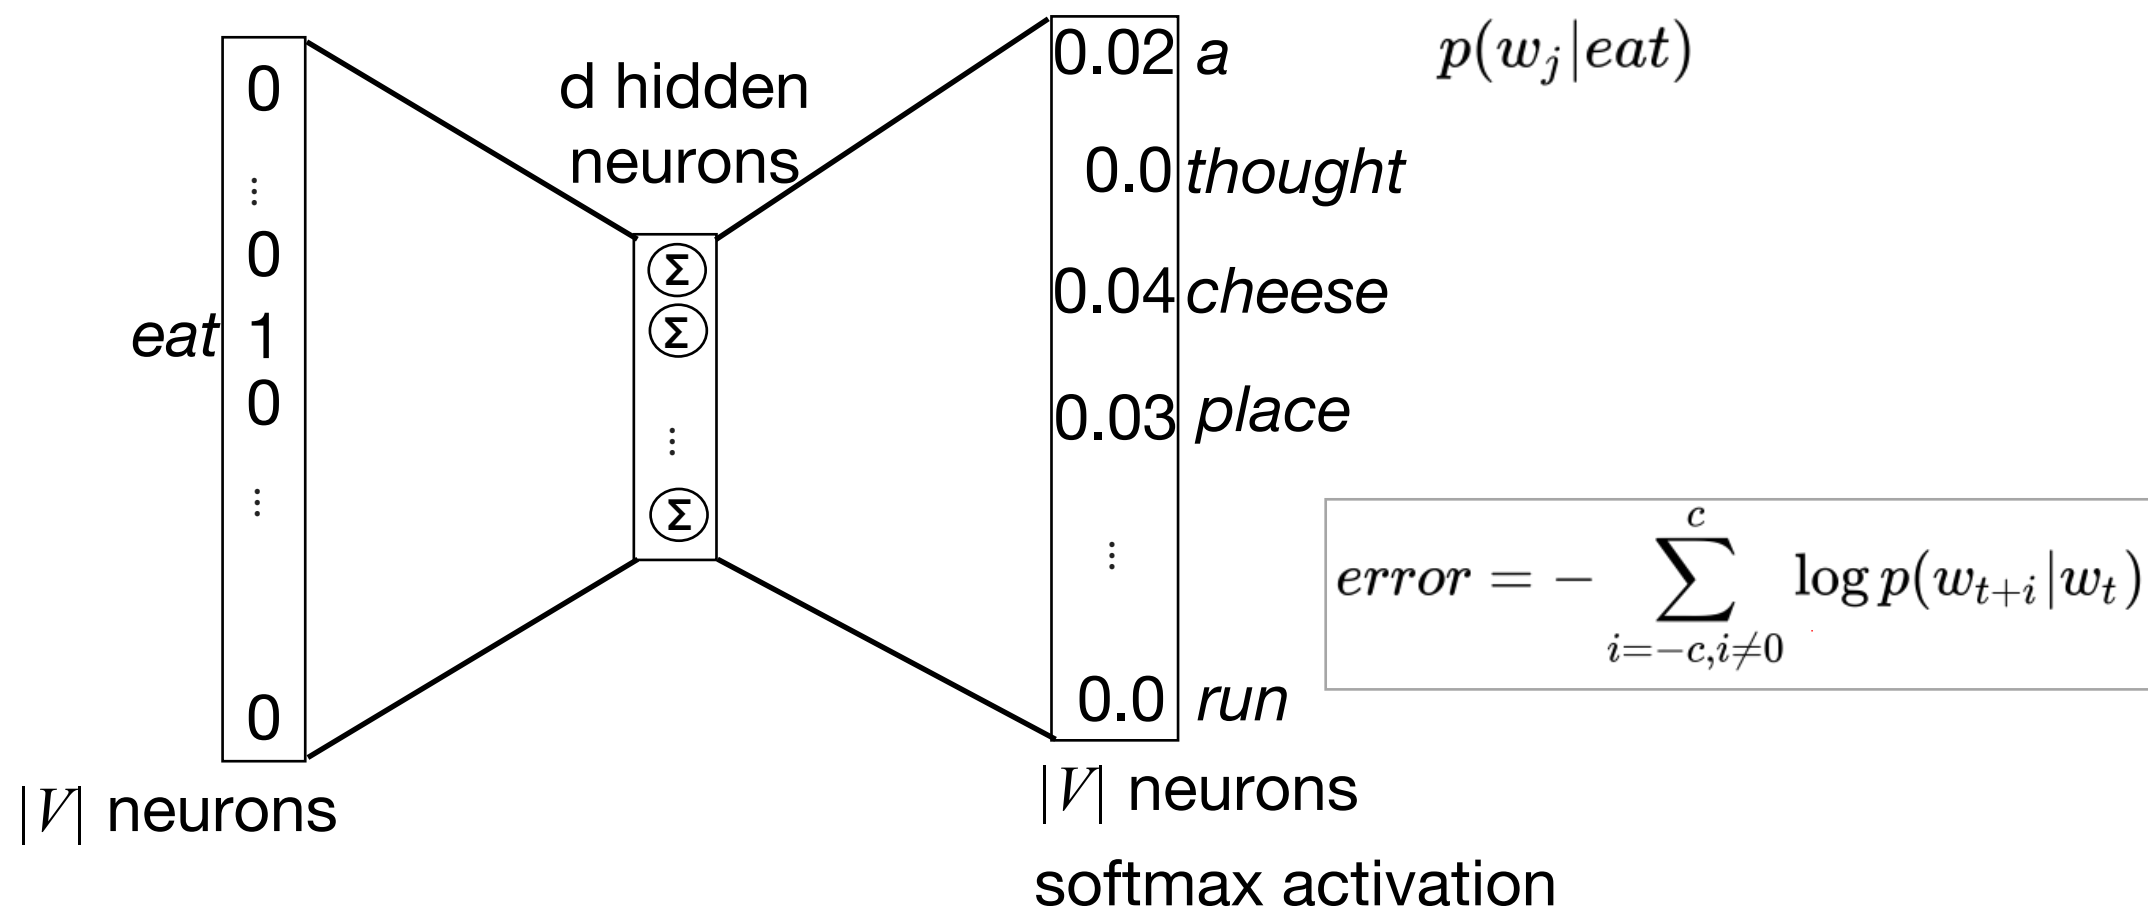
$$\frac{\partial \text{Loss}}{\partial w} = \left(\frac{\partial \text{Loss}}{\partial f(x)} \right) \left(\frac{\partial f(x)}{\partial w} \right)$$

**to compute these
we have to know
the derivate of the
function f**

Autoencoders Embeddings (Word level semantics)

Skip-Gram Model

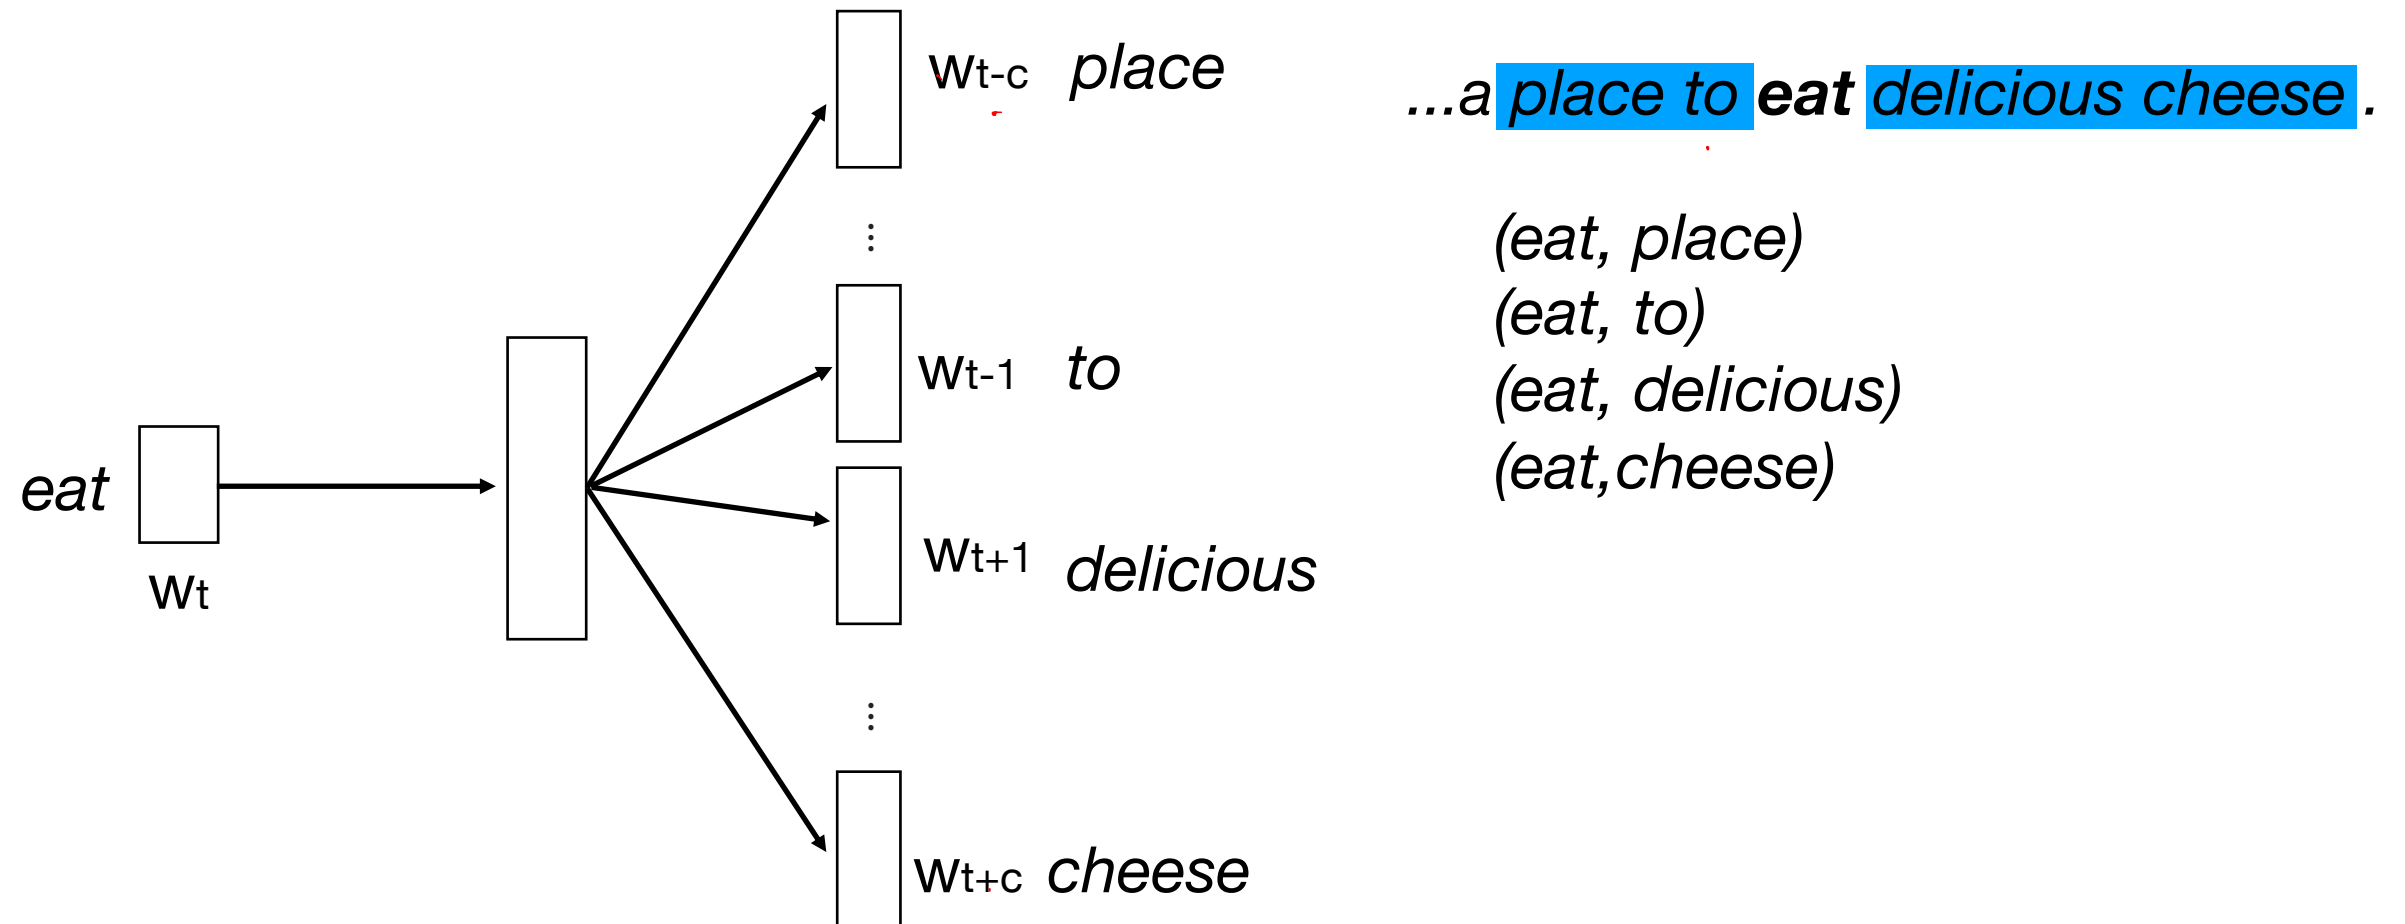
- Input:
A single word in one-hot representation.
- Output: probability to see any single word as a context word.



- Softmax function normalizes the activation of the output neurons to sum up to 1.0.

Skip-Gram Model


- Compute error with respect to each context word.



- Combine errors for each word, then use combined error to update weights using back-propagation.

$$error = - \sum_{i=-c, i \neq 0}^c \log p(w_{t+i} | w_t)$$

Continuous Bag-of-Words Model (CBOW)


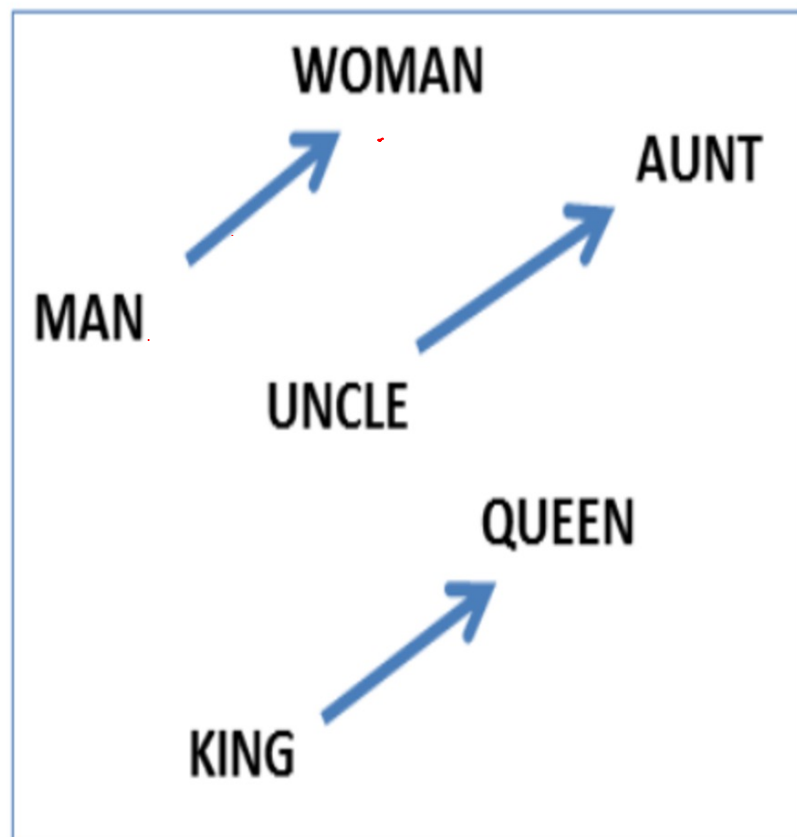


- Input: Context words. Averaged in the hidden layer.
- Output: Probability that each word is the target word.

Embeddings are Magic

(Mikolov 2016)

$\text{vector}(\text{'king'}) - \text{vector}(\text{'man'}) + \text{vector}(\text{'woman'}) \approx \text{vector}(\text{'queen'})$



Application: Word Pair Relationships


Table 8: *Examples of the word pair relationships, using the best word vectors from Table 4 (Skip-gram model trained on 783M words with 300 dimensionality).*

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

Using Word Embeddings

- Word2Vec:
 - <https://code.google.com/archive/p/word2vec/>
- GloVe: Global Vectors for Word Representation
 - <https://nlp.stanford.edu/projects/glove/>
- Can either use pre-trained word embeddings or train them on a large corpus.

Word embeddings

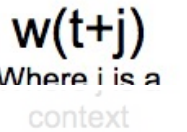


$$error = - \sum_{i=-c, i \neq 0}^c \log p(w_{t+i} | w_t)$$

softmax activation

- pas

- ## Reference



$w(t-2)$
Or, two spots behind
 $w(t)$

$w(t-1)$
Or, one spot behind
 $w(t)$

 $w(t+1)$
or, one spot ahead
 $w(t)$ $w(t+2)$
Or, two spots ahead
 $w(t)$

How can we build a sentence representation using word-level distributional representations?

Acknowledgments

- Some slides by Chris Kedzie